

# Neural Networks Deep Learning and friends...



# Agenda

- AI – Where does DL fit?
- How is it different from “Classical” Programming
- Path from Design to Production
- What is a Neuron
  - With live examples
- Neuronal Network:  
TensorFlow playground
- Machine Learning
  - Layers
  - See it at work on TF Playground
- Training a Model
  - Building the Network
- Car Insurance Demo
- Character Recognition
  - With OpenCV

# The Tools

- Java
- Python
- Jupyter Notebooks
  - For Python
  - For Java
- TensorFlow
- Keras
- OpenCV

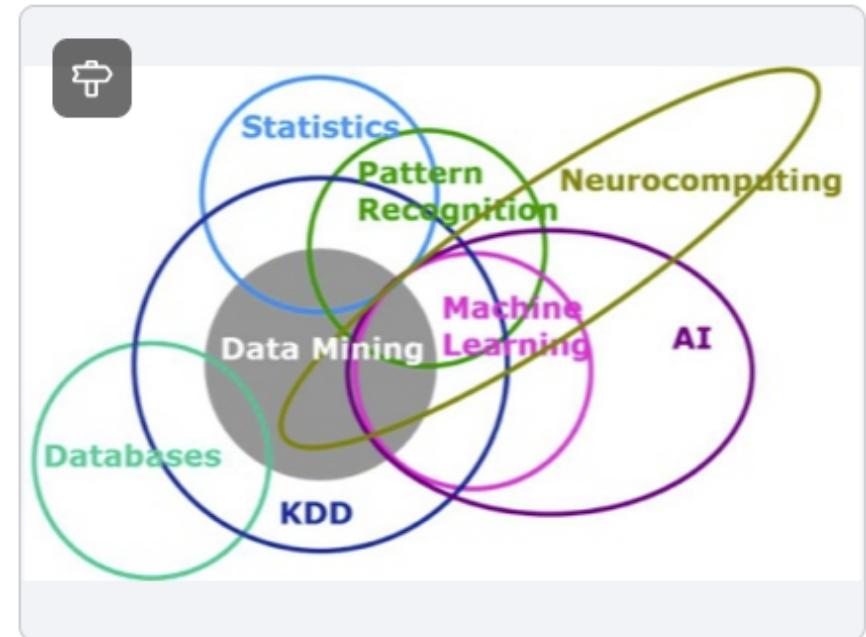
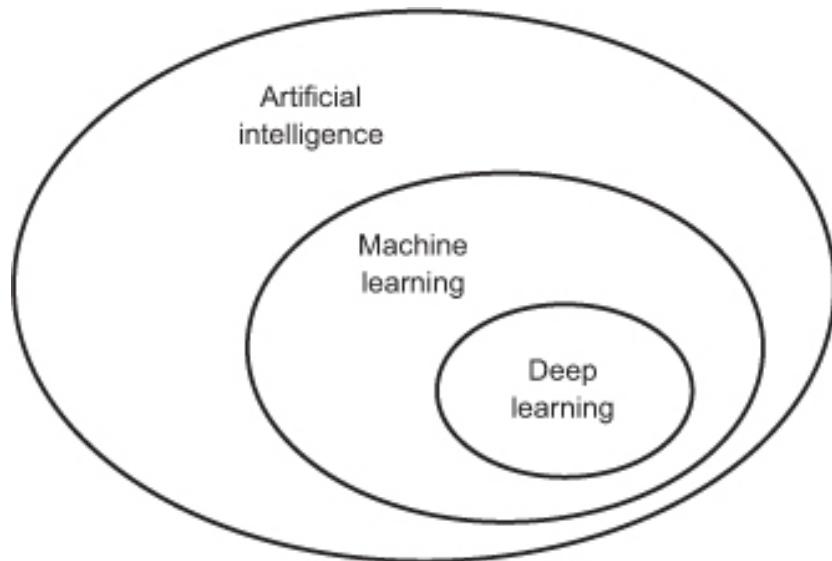
# We will

- See what ML/DL is, why it is useful, and what challenges it might face
- See what a neuron is, where and how it fits
- See what a Neuron Network is and how to build one
- See how to use a Neuron Network to train a model
- Use a simple example, train a model, and use it
- Take a more complex (aka real) example, end to end

# Machine Learning

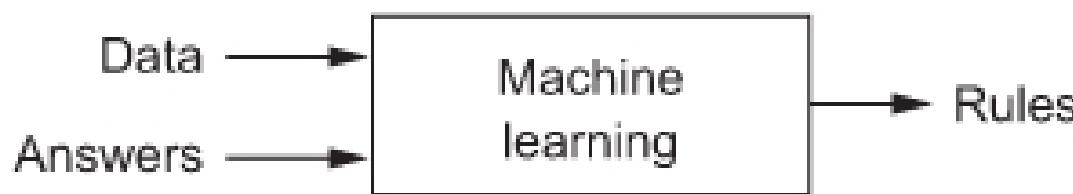


# Machine Learning



(KDD: Knowledge Discovery in Databases)

# How is that different?



For the learning phase – aka model training – you need a lot of data (the more the better), separated in two groups:

- Input data, parameters of the equation (like age, speed, mileage)
- Known output, result (risk level)

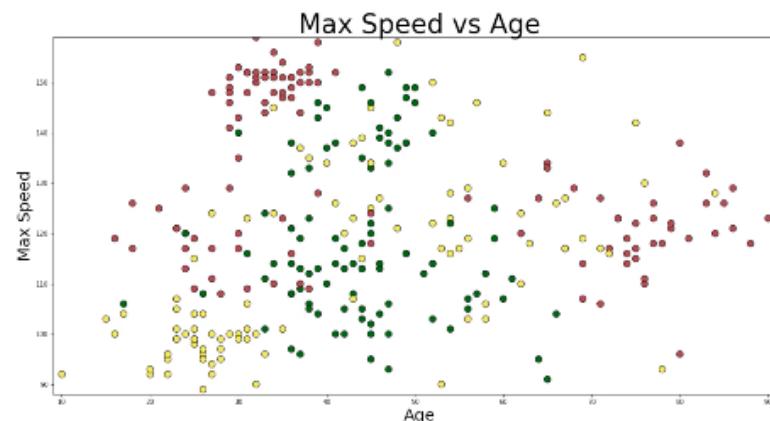
The *learning phase* is the process that tries to find the way to reach the expected result with the input parameters **for all the data** of the *training data set*.

# Where it fits

- It is all about the **MODEL**
- The **MODEL** is the entity that will eventually apply the rules found during the *Training*
  - *Training* the **MODEL** is equivalent to *finding the Rules* to apply to future input data to *predict* the output
- Once *trained*, the model is deployed to *production*
  - Real world and real time data can be sent to the **MODEL** to get its *Predictions*
- A deployed **MODEL** can be re-trained and re-deployed to *production*

# Our first example

- Let's say you are an insurance company, and you want to be able to predict the risk associated with a potential customer, based on
  - His age
  - The maximum speed of his car
  - The mileage he drives per year
- We do have existing data, from some insurance company database



# Apply a predefined rule

```
red = 0
green = 1
yellow = 2

def evaluate_risk(age, speed, miles_per_year):
    if age < 25:
        if speed > 140:
            return red # Crazy young guy, car too fast
        else:
            return yellow # Car is slow enough for medium risk

    if age > 75:
        return red # Get off the road, old man!

    if miles_per_year > 30:
        return red # You drive too much

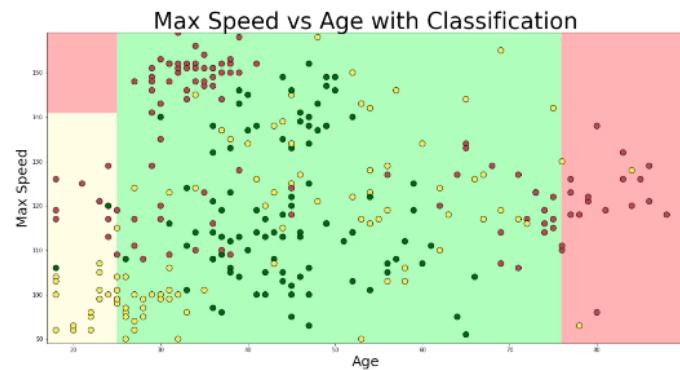
    if miles_per_year > 20:
        return yellow

    return green # otherwise, low risk
```

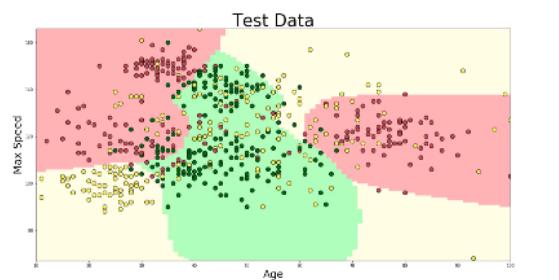
Notebook here: <http://localhost:8888/notebooks/NeuralNetworks.ipynb>

# What we have vs what we want

- This would produce something like this:



- As opposed to what we would expect:



# Not good, obviously!

- We need another way to tackle that.
- Let's try to teach the machine ...



# Neurons

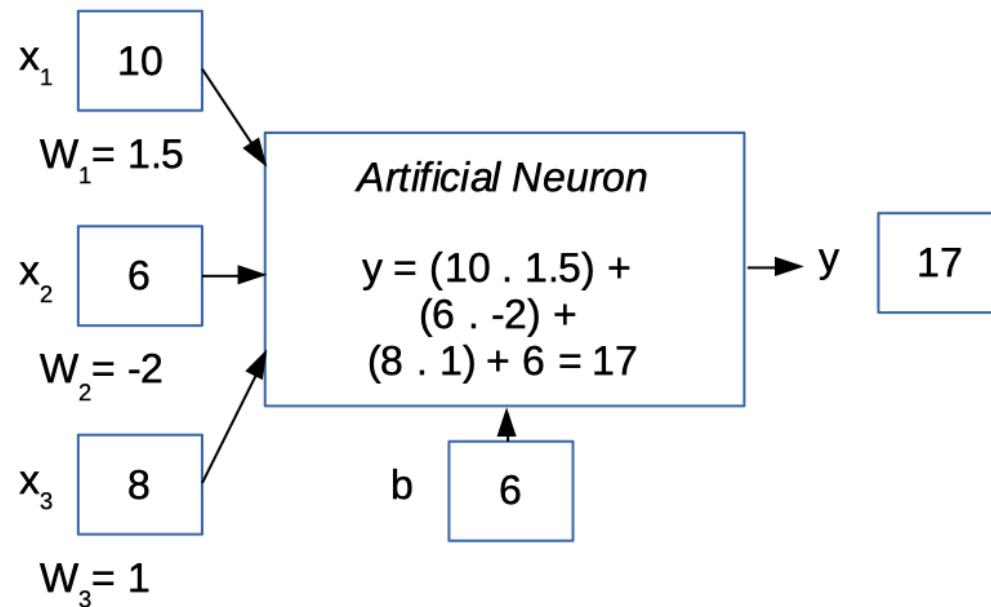
(aka Artificial Neurons)



# What's a Neuron

- First of all, ***this is not new!***
- It looks a lot like a mathematical function, because this is what it is
- A Neuron has a single numerical output (usually called **y**)
- It can have several inputs (also called dimensions), usually named **X<sub>i</sub>**, associated with a weight **W<sub>i</sub>**.
- It also has a constant, called the bias (noted **b**)
- Its formula is
$$y = \left( \sum_{i=1}^n x_i \cdot W_i \right) + b$$
- In the next sections, they will be part of the *hidden layers*

# Example



- Yes, it is very simple!
- [Java Notebook](#)
- [Web page](#)

# Neuron Networks

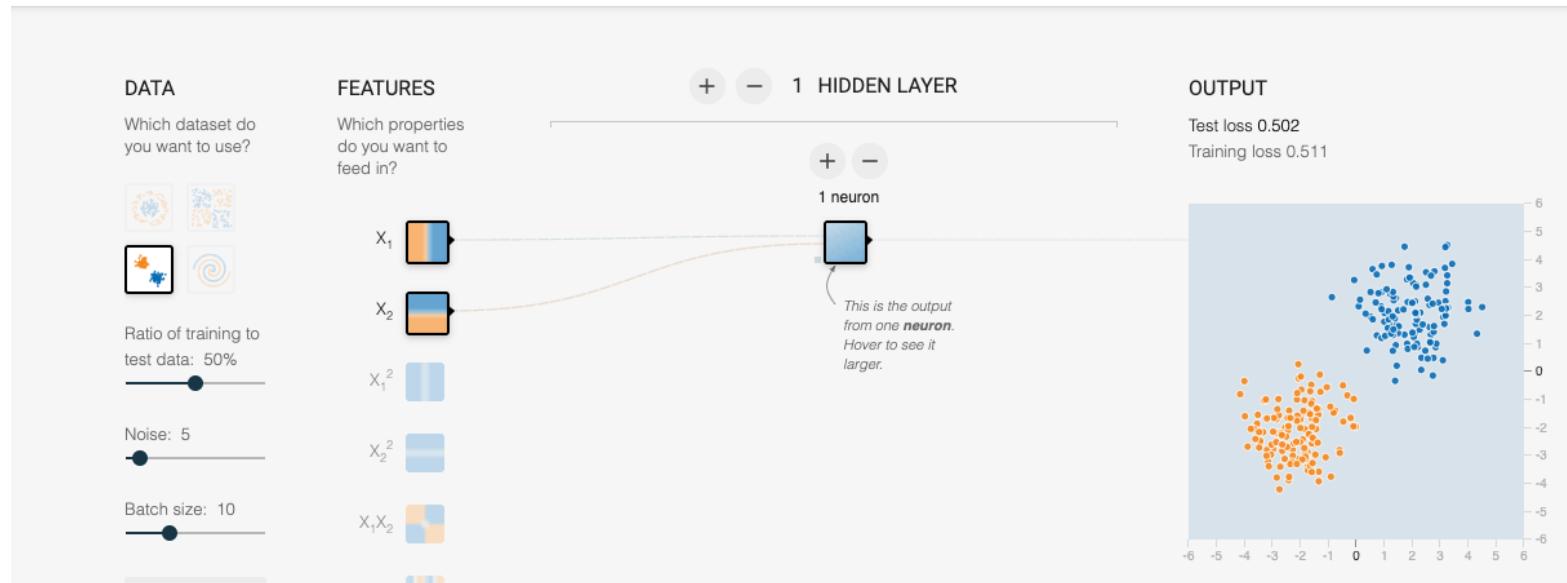


# Neural Networks

- Introducing [TensorFlow Playground](#)
  - The different datasets
  - See the input nodes
  - The ***hidden layers***
    - **Note:** The number of hidden layers is called the ***depth*** of the network. This is where the ***Deep*** in ***Deep-Learning*** comes from.
    - The weights (lines thickness)
    - The bias (small squares at the bottom left of the neurons)
    - The run button
      - Notice the loss graph
        - Test, training losses
    - Learning rate
    - Activation
    - ...

Note: A ***Tensor*** is the vector containing the input data

# TensorFlow Playground



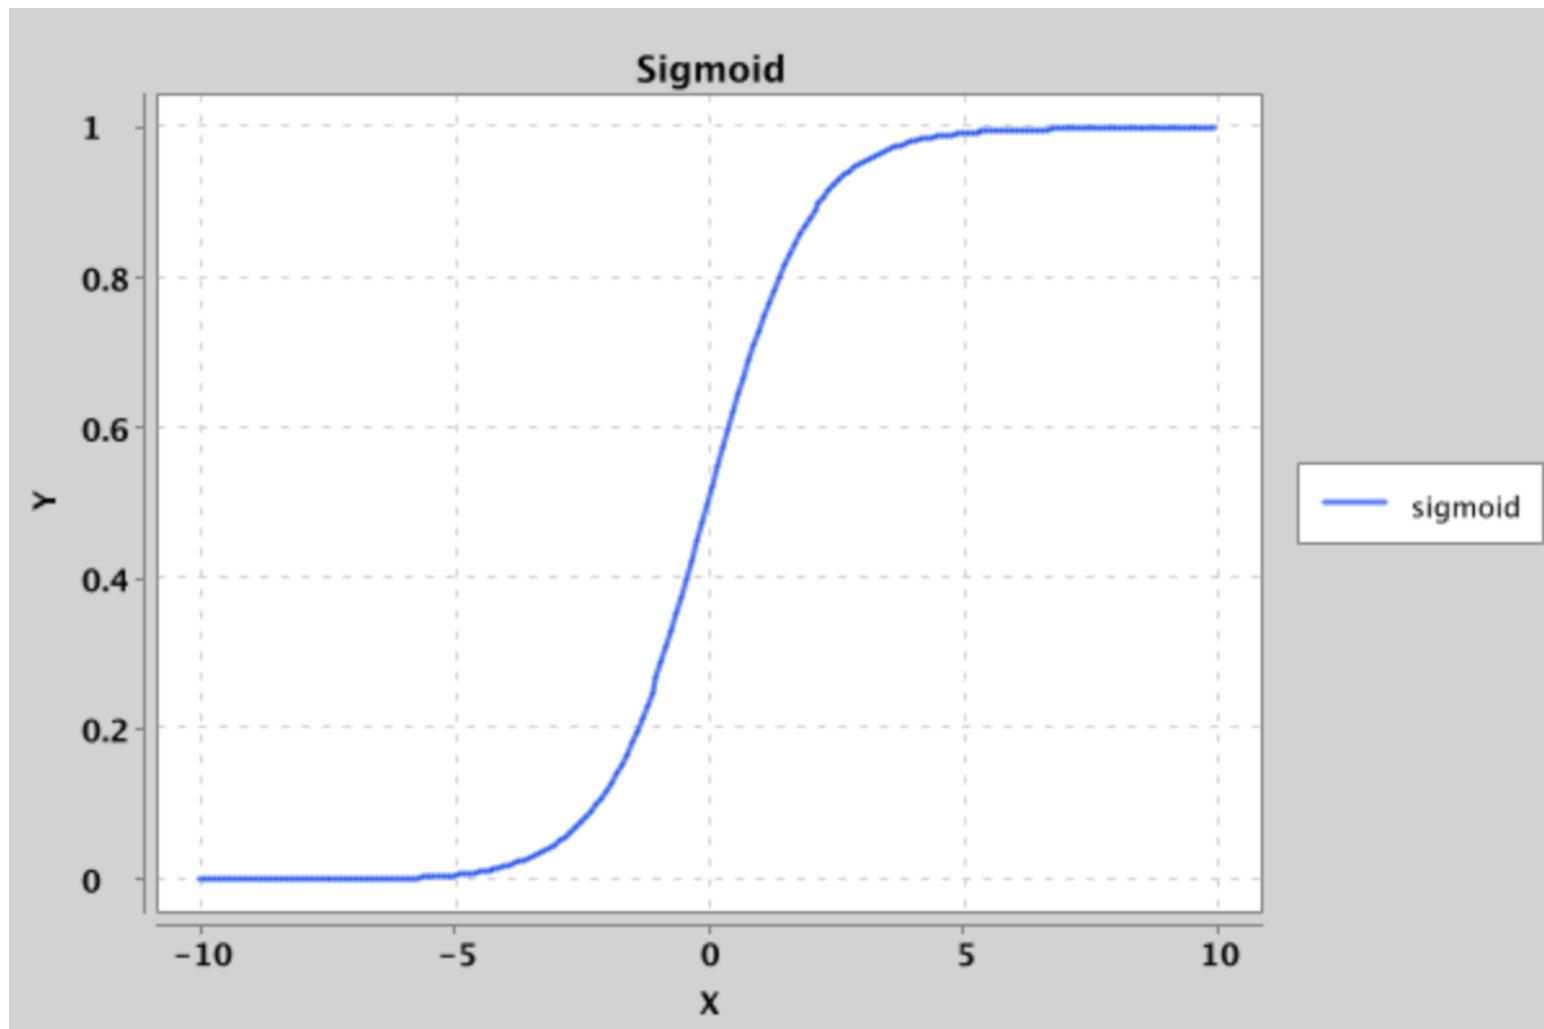
- Here we have 2 input dimensions,  $x_1$  (abscissa) and  $x_2$  (ordinate).
- The **weight** of each dimension is represented by the thickness of the line between the dimension and the neuron(s)
- The **bias** of the neuron is represented by the (very) small little gray square at the bottom left of the neuron(s).
- *The output will be a single number, ranging from -1 to +1.*
- The background of the graphic on the right will reflect the prediction, it will be orange if  $y < 0$ , and blue if  $y > 0$ .
- You can change the weights associated with each box (dimensions and neuron(s)), as well as the bias of each neuron.

# Activation Functions

- An activation function is triggered after weight(s) and bias are applied
- It “re-centers” the result ( $y$ ) of the neuron
  - In the previous TF example, the result will be “orange” or “blue”. A value like -1 or +1, 0 or +1 is absolutely good enough
- Common activation functions are
  - Sigmoid
  - TanH
  - Step
  - ReLU
  - SoftMax
- See them live [here](#)

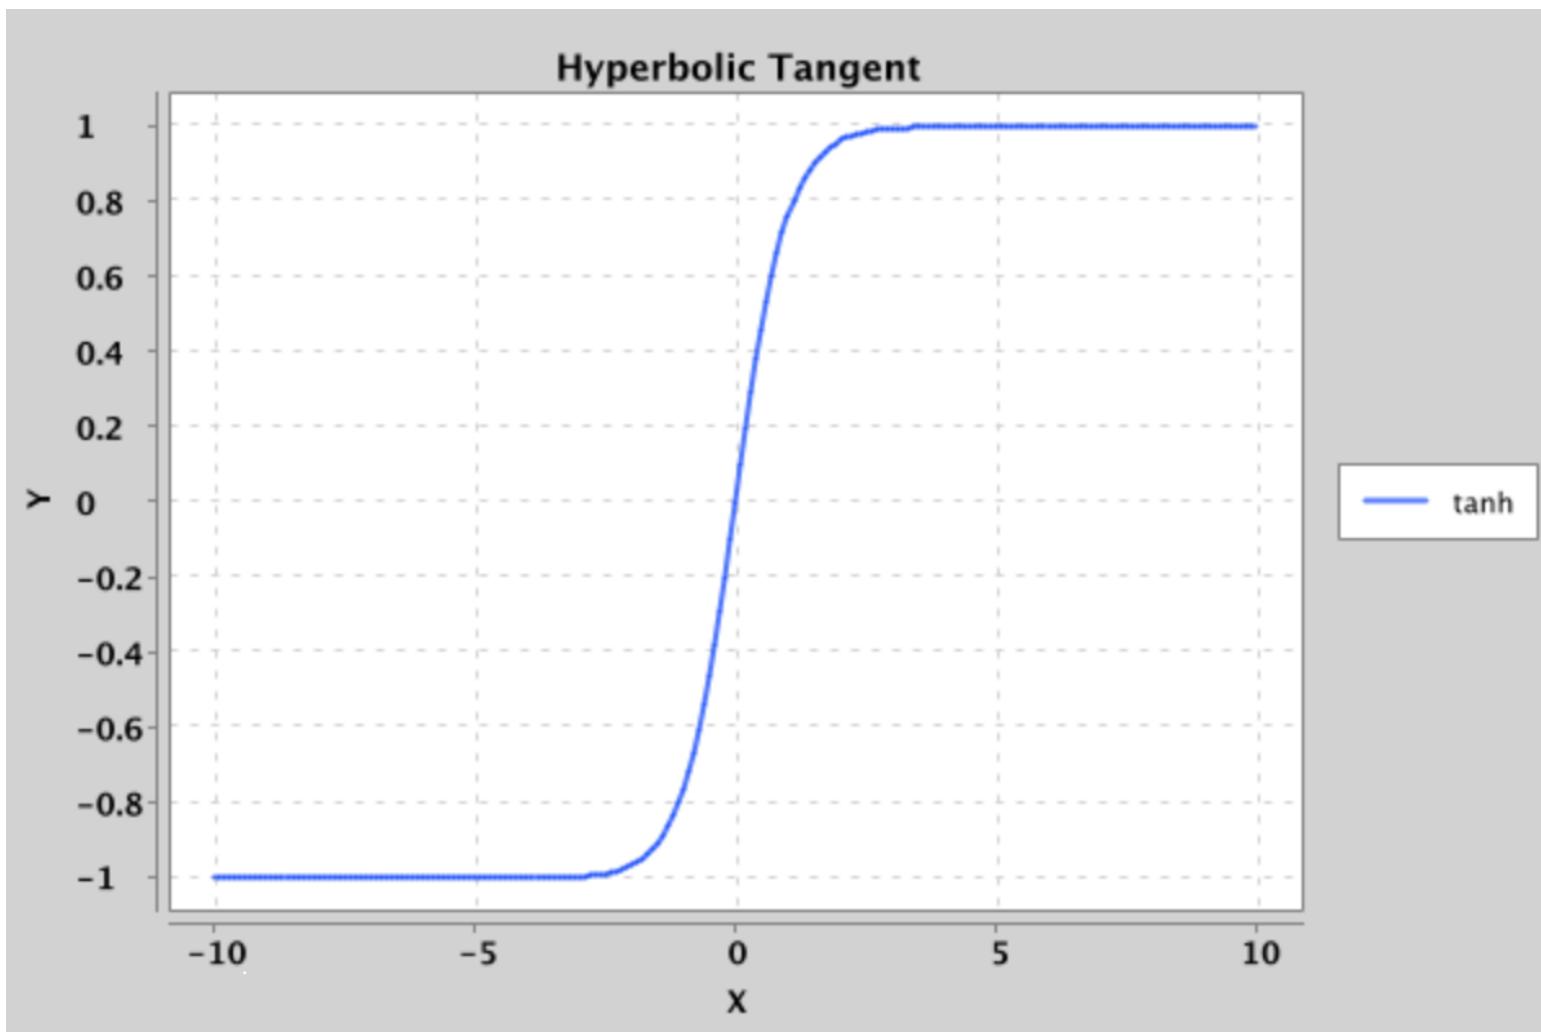
# Sigmoid

$$x \rightarrow (1 / (1 + e^{-x}))$$



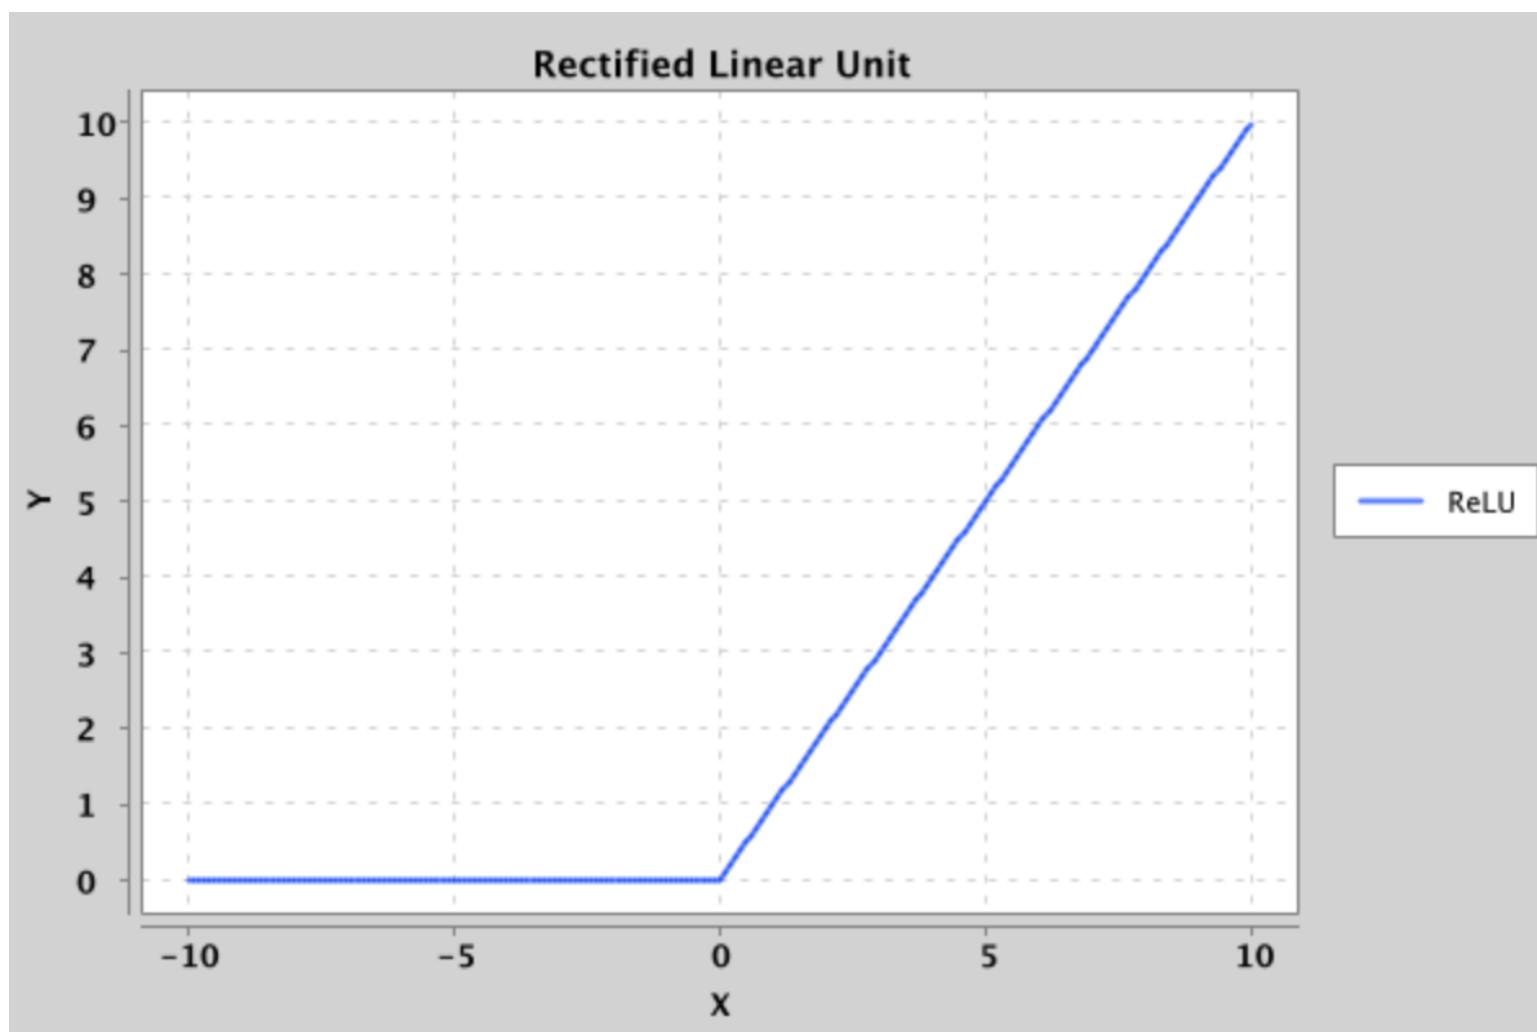
# Hyperbolic Tangent

$$x \rightarrow (e^x - e^{-x}) / (e^x + e^{-x})$$



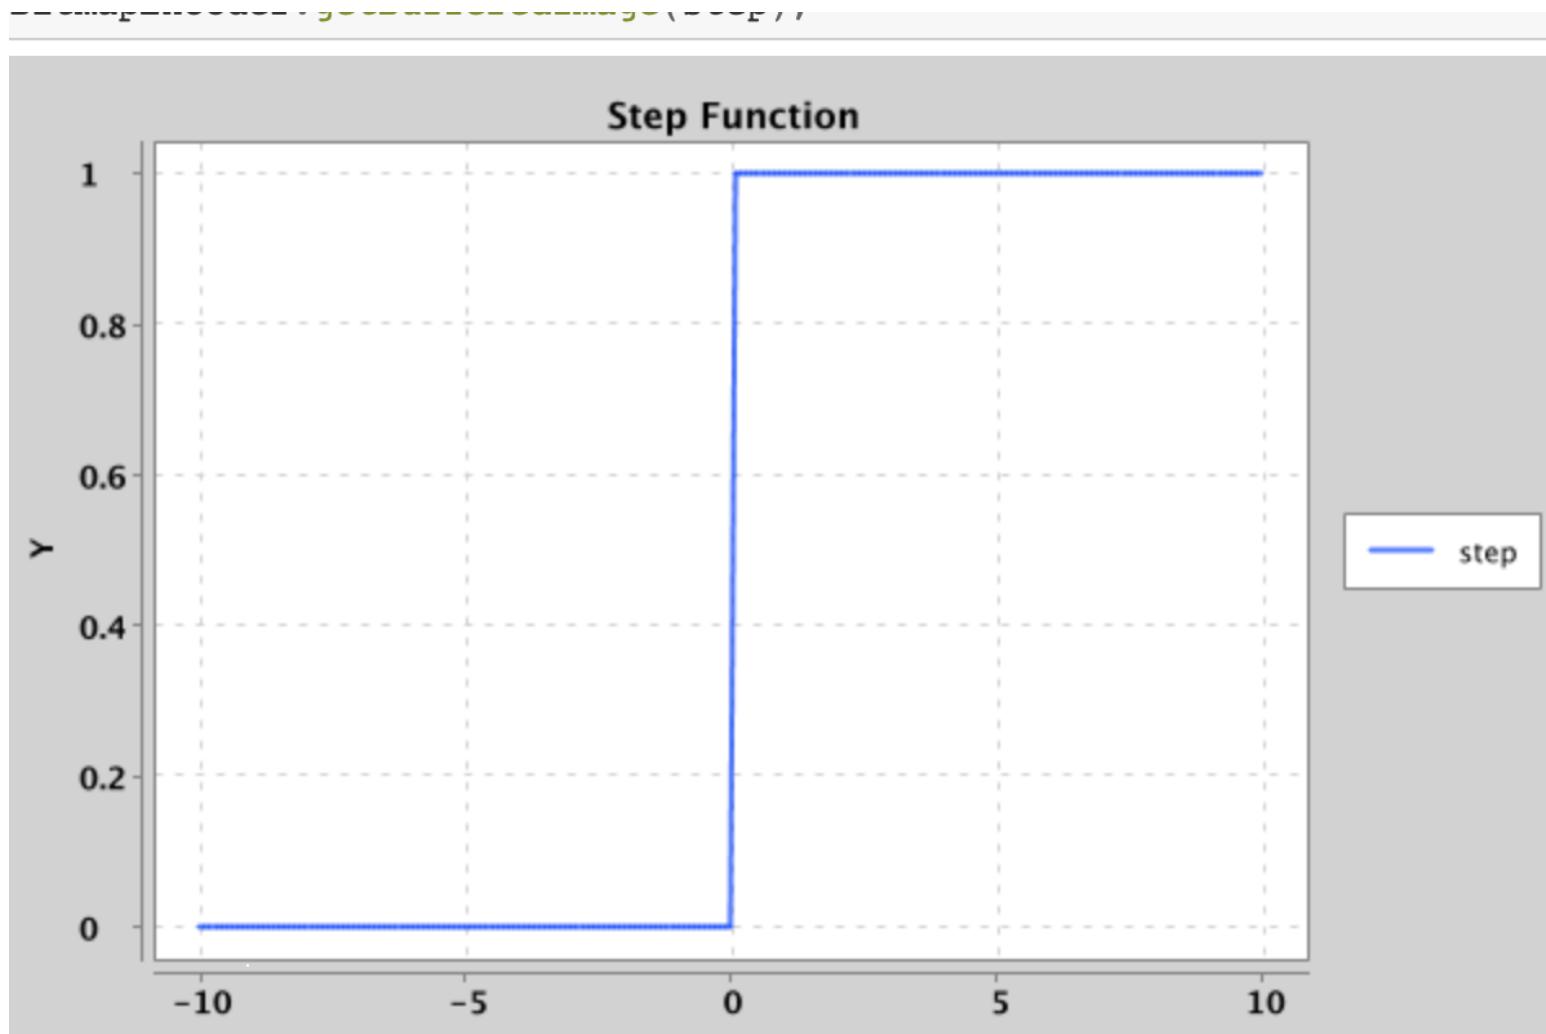
# Rectified Linear Unit (ReLU)

$$x \rightarrow x < 0 ? 0 : x$$



# Step

$x \rightarrow x < 0 ? 0 : 1$



# SoftMax

- Depends on the number of expected output values, suited for Classification problems
- $\sum_{i=1}^n \text{value}[i] = 1$

For example, in the car insurance case

We say:

- Red = 0
- Yellow = 1
- Green = 2

A result like [0.01, 0.95, 0.04]

is red=1%, yellow=95%, green=4%

# Loss

- That is **THE** thing leading to an accurate model
  - The loss is the difference between the reality and what the model comes up with
  - Finding the minimal loss ***is the goal*** of the training
  - The smallest the loss, the best the model
    - Obviously!
- Loss can be set for
  - Training data
  - Test data

# Demo

# TensorFlow Playground



TensorFlow Playground

# How does the Network learn?



# Minimizing the Loss

- For each tuning position of the network (weights, bias, for *all* neurons in *all* layers)
- Compare the calculated result with the expected one
- Use Mean Square Error (MSE) to magnify the loss
- Plot it against epoch
- Find the spot where loss is the smallest
- See [here](#)

*Mean Squared Error*

$$MSE = \frac{1}{n} \sum_{i=1}^n (Y_i - \hat{Y}_i)^2$$

# Training a TensorFlow model



# Let's see some code

## Car Insurance Demo:

- Jupyter notebook, [training a model](#)
  - Setting up the network requires some... intuition
- Convert model to TensorFlowJS
  - `tensorflowjs_converter --input_format keras  
./insurance.h5 ./tfjs`
- Using the model in “Production” ([Web UI](#))

Closer to the real world



# Hand-written character recognition

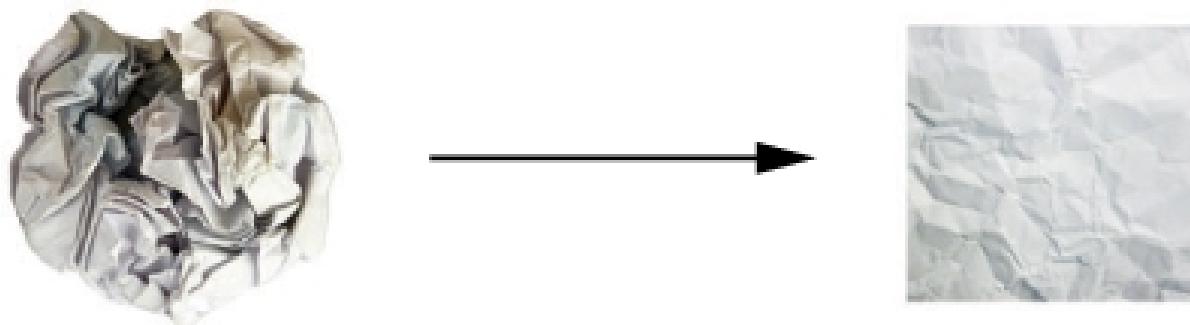
- Now we've seen the bits and pieces, let's put them at work
- We are going to ***train*** a model to recognize hand-written digits, in [0..9]
- We will use a webcam to “see” a notepad
- Take a snapshot, ***prepare*** the data returned by the camera
- Send them to the ***model*** to ***predict*** what digit it is
- Let's go

# Hand-written character recognition

- We need to “see” the user writing on the notepad
- Introducing **OpenCV**

# Hand-written character recognition

- A word about data preparation
- It is like unraveling the data



# Hand-written character recognition

- The model will be trained with B&W images, 28x28 pixels
- This is the **mnist** data set, that comes with Keras

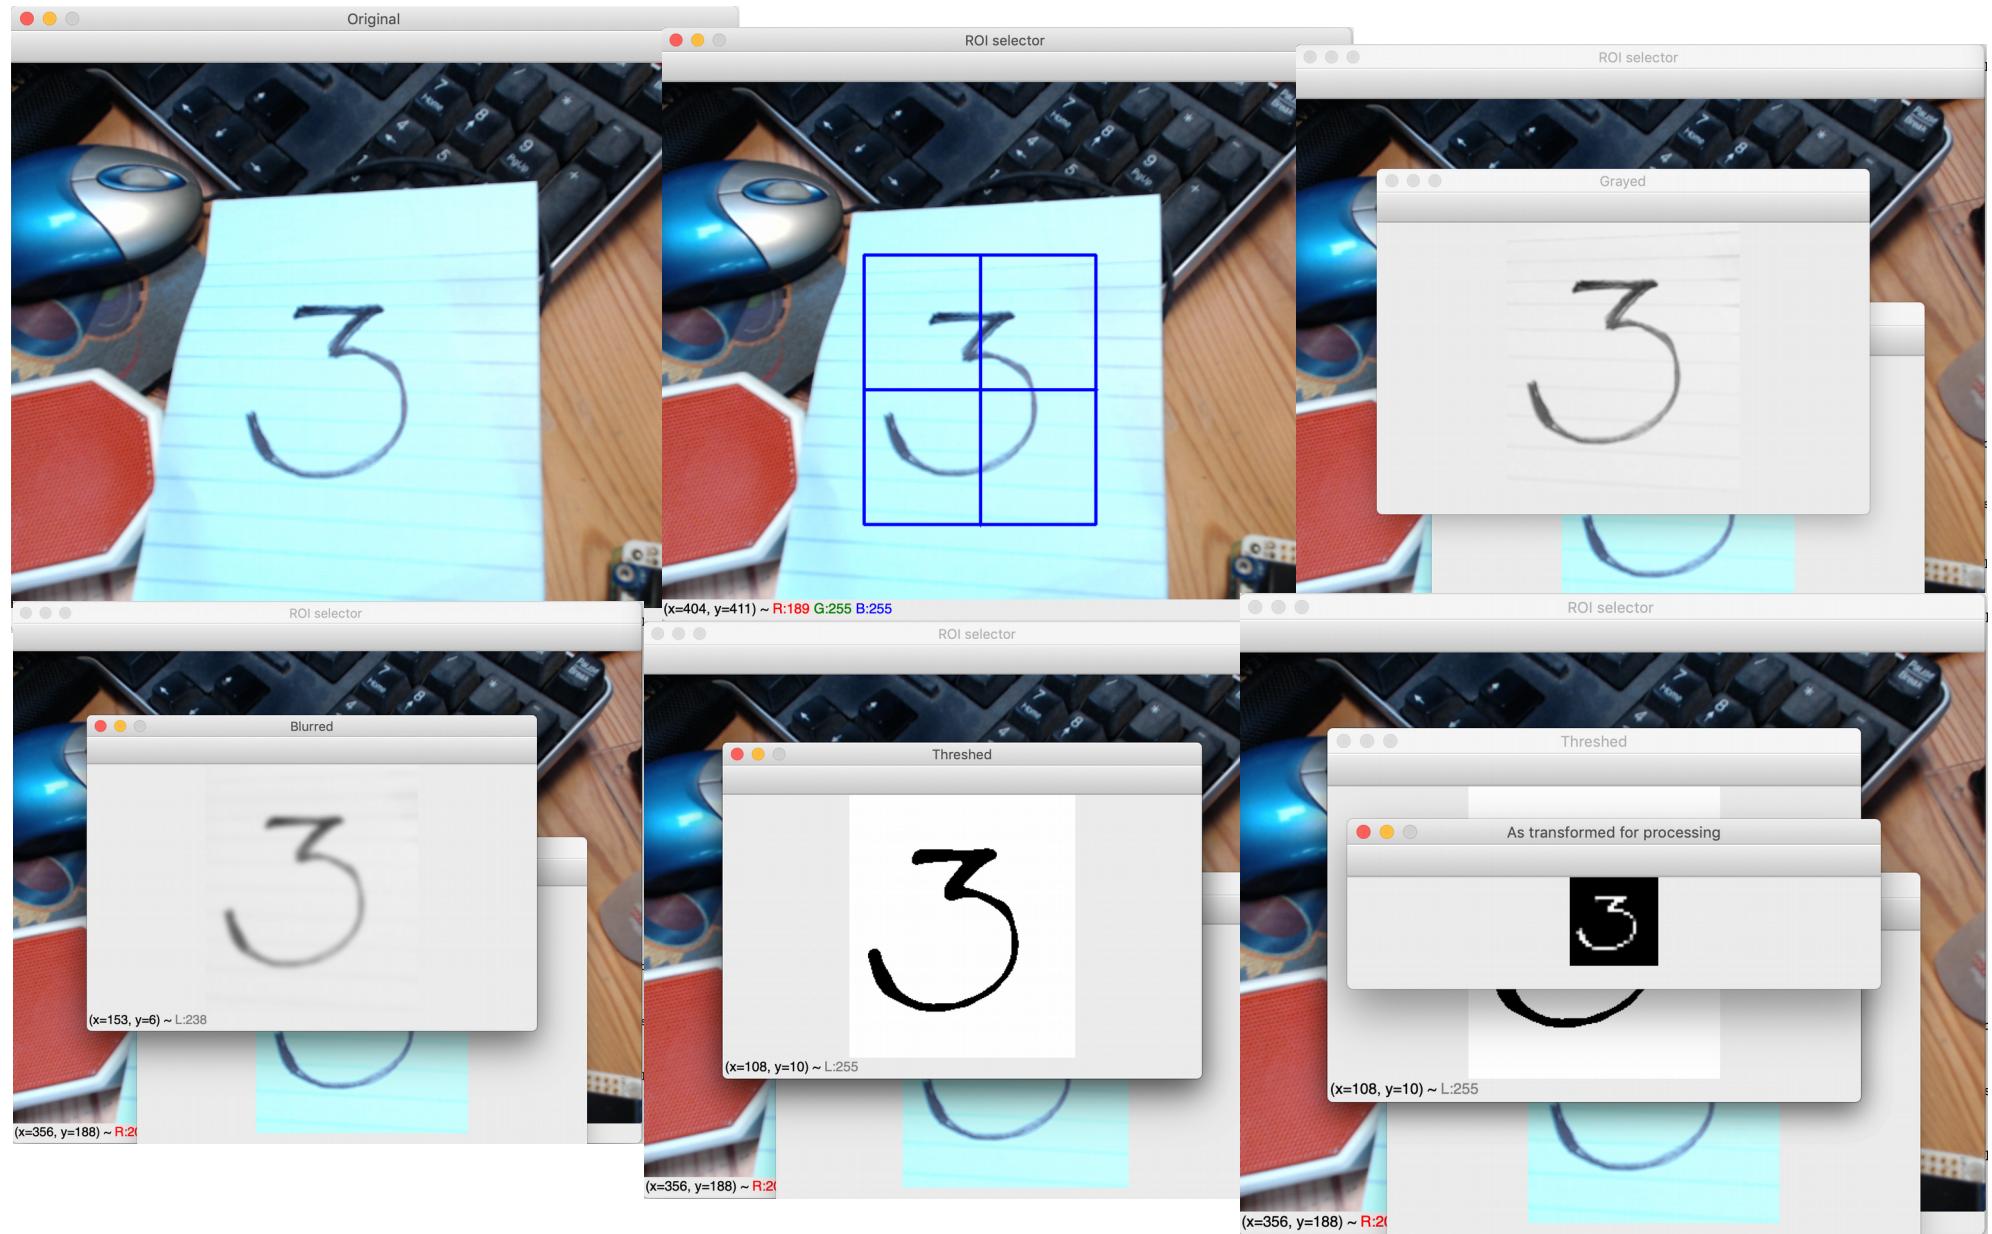


8	2	4	3	3	6	9
8	2	4	3	3	6	9
9	5	8	6	7	0	6
9	5	8	6	7	0	6
2	6	3	9	1	7	4
2	6	3	9	1	7	4
8	8	9	0	3	0	5
8	8	9	0	3	0	5
2	9	4	1	0	3	7
2	9	4	1	0	3	7
5	8	7	7	8	2	9
5	8	7	7	8	2	9
5	5	1	2	6	4	2
5	5	1	2	6	4	2

# Hand-written character recognition

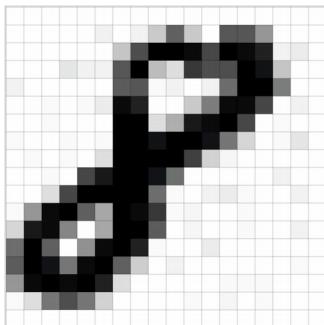
- Using OpenCV, we will
  - Take a snapshot
  - Select a Region Of Interest (ROI)
  - Gray it (no color)
  - Blur it to get rid of potential background
  - Thresh it (minimize number of colors)
  - Resize it as expected by the model

# Hand-written character recognition



# Hand-written character recognition

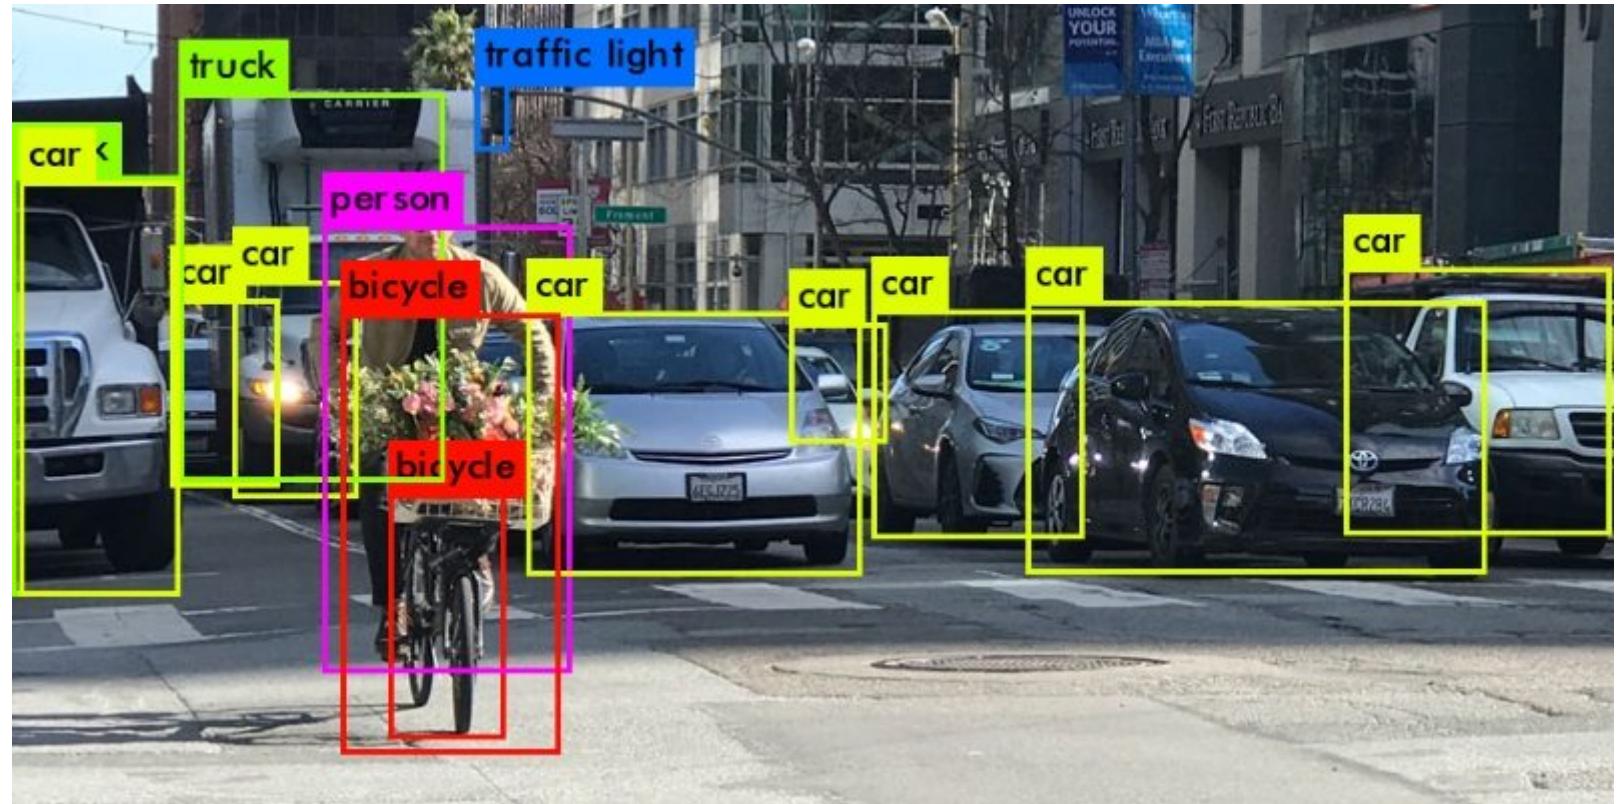
- Data preparation



# Hand-Written Digits Demo



# Even closer to the real world: Self driving cars



- Detect objects around
  - cars, bikes, pedestrians, traffic lights, curbs, etc
- Classify them, and make the right decision
  - brake, stop, slow down, turn right, etc

# Final Bonus.

59 seconds



# Recommended reading

- Manning's Deep Learning Crash Course
  - by Oliver Zeigermann
- Manning's Grokking Deep Learning
  - by Andrew Trask
- Manning's Deep Learning with Python
  - by François Chollet
- All the material on github.

